

Mapping Aggregates to Azure CosmosDB



Julie Lerman

EF Core Expert and Serial DDD Advocate

@julielerman | thedatafarm.com

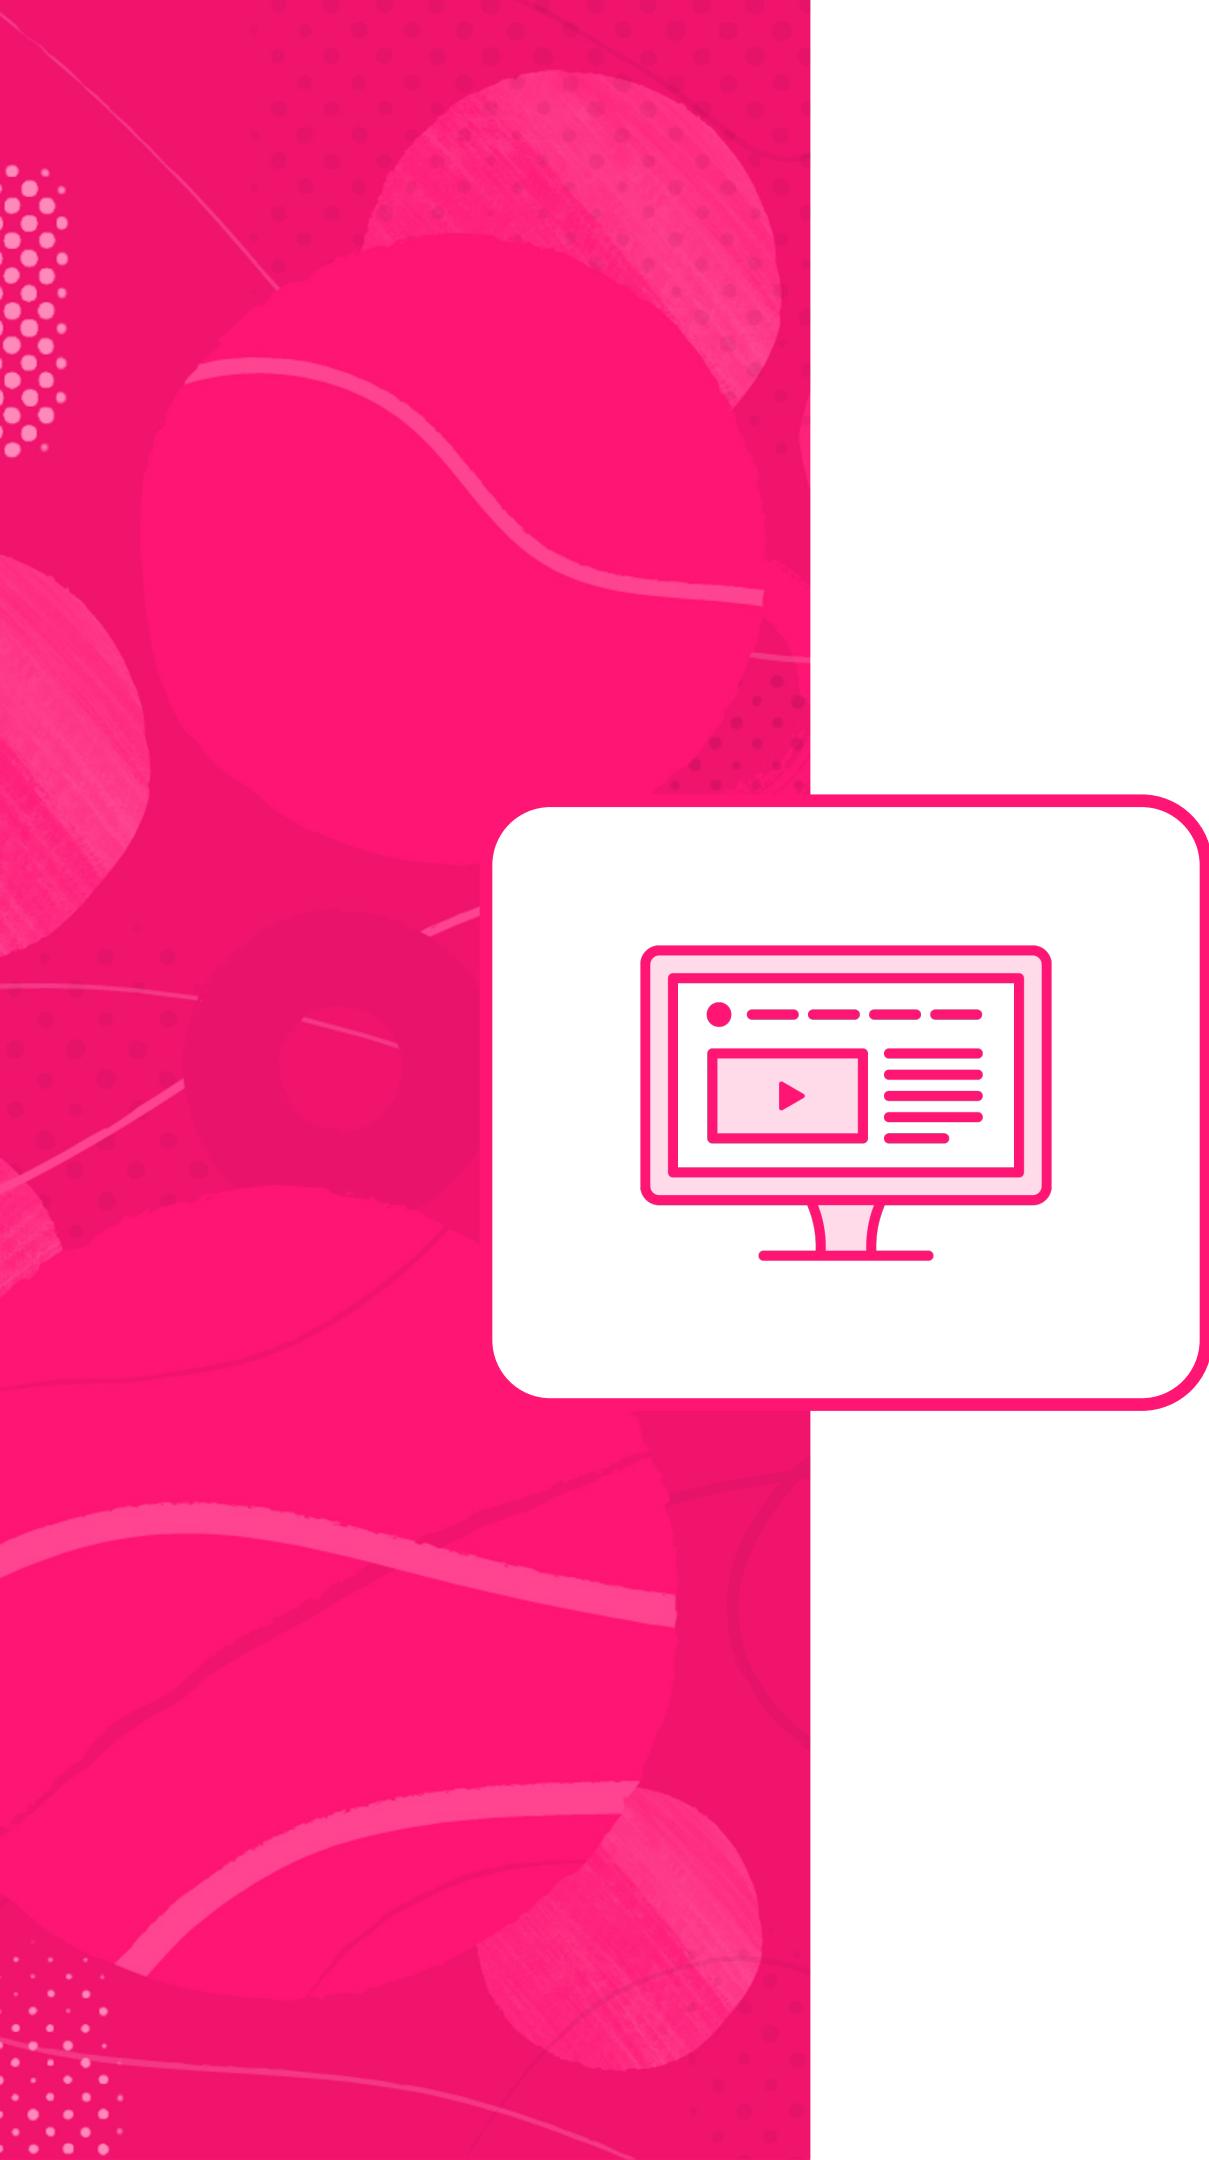


Why Use an ORM* With Unstructured Data?

Leverage existing skills rather than have to learn a new set of APIs for interacting with CosmosDB

* EF Core is much more than just an ORM ☺





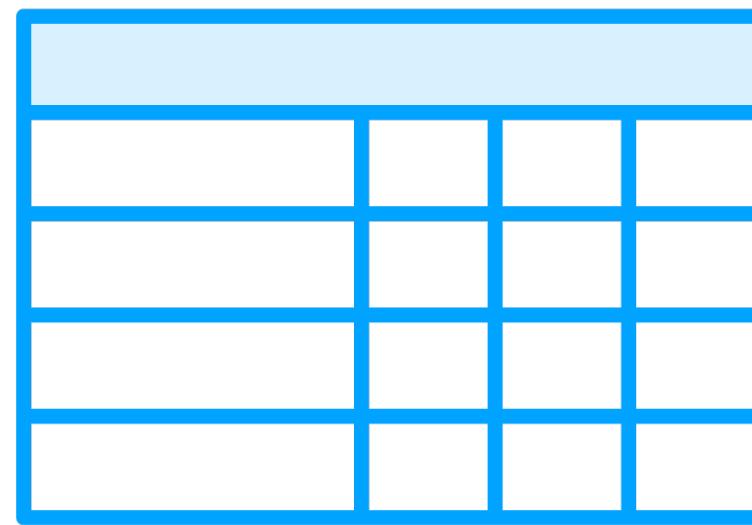
Much More About this Provider

**Using EF Core 6 with Azure
Cosmos DB**

Jurgen Kevelaers



Document Database Removes Impedence Mismatch



**Relational database stores data
in a highly structured schema**

{JSON}

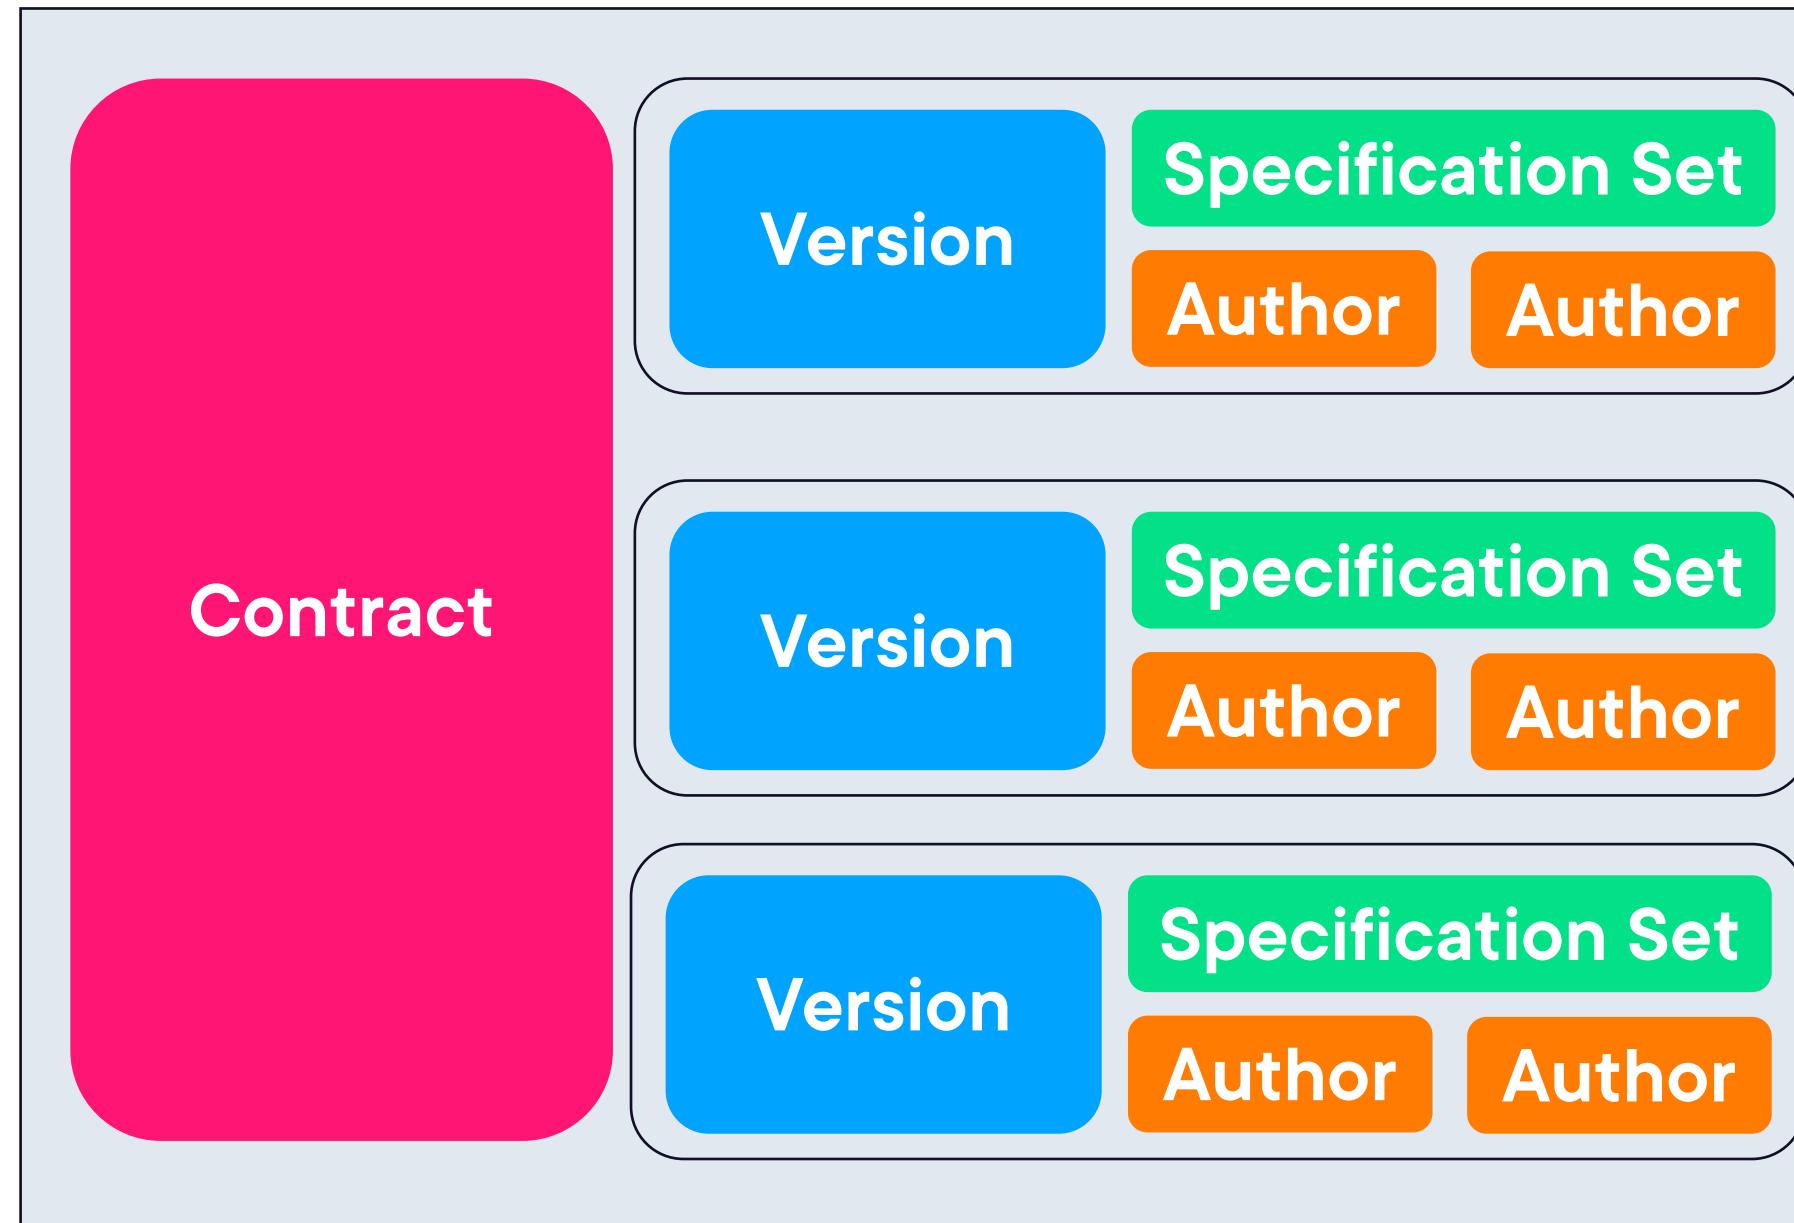
**Document database stores data
in JSON documents as
unstructured data**



**This module will focus on
the differences that are
relevant to mapping the
DDD-guided aggregate.**



CosmosDB Can Store the Aggregate in One Document



```
{  
  "ContractId": "123",  
  "ContractNumber": "xyz",  
  "Versions": [  
    {  
      "VersionId": "456",  
      "Accepted": false,  
      "Authors": [  
        {  
          "Email": "a@a.com",  
          "Name": {  
            "FirstName": "Julie"  
            "LastName": "Lerman"  
          },  
        },  
        {"Email": "b@b.com", "Name": { "FirstName": "John", "LastName": "Doe" }}  
      ],  
      "Specs": {  
        "AdvanceAmountUSD": 1,  
        "DigitalRoyaltyPct": 10  
      }  
    }  
  ]  
}
```



Some Key Mapping Differences

Some mappings don't apply
to CosmosDB

Automatically embedded
value objects and related data

Store collections and dictionaries
of scalar data

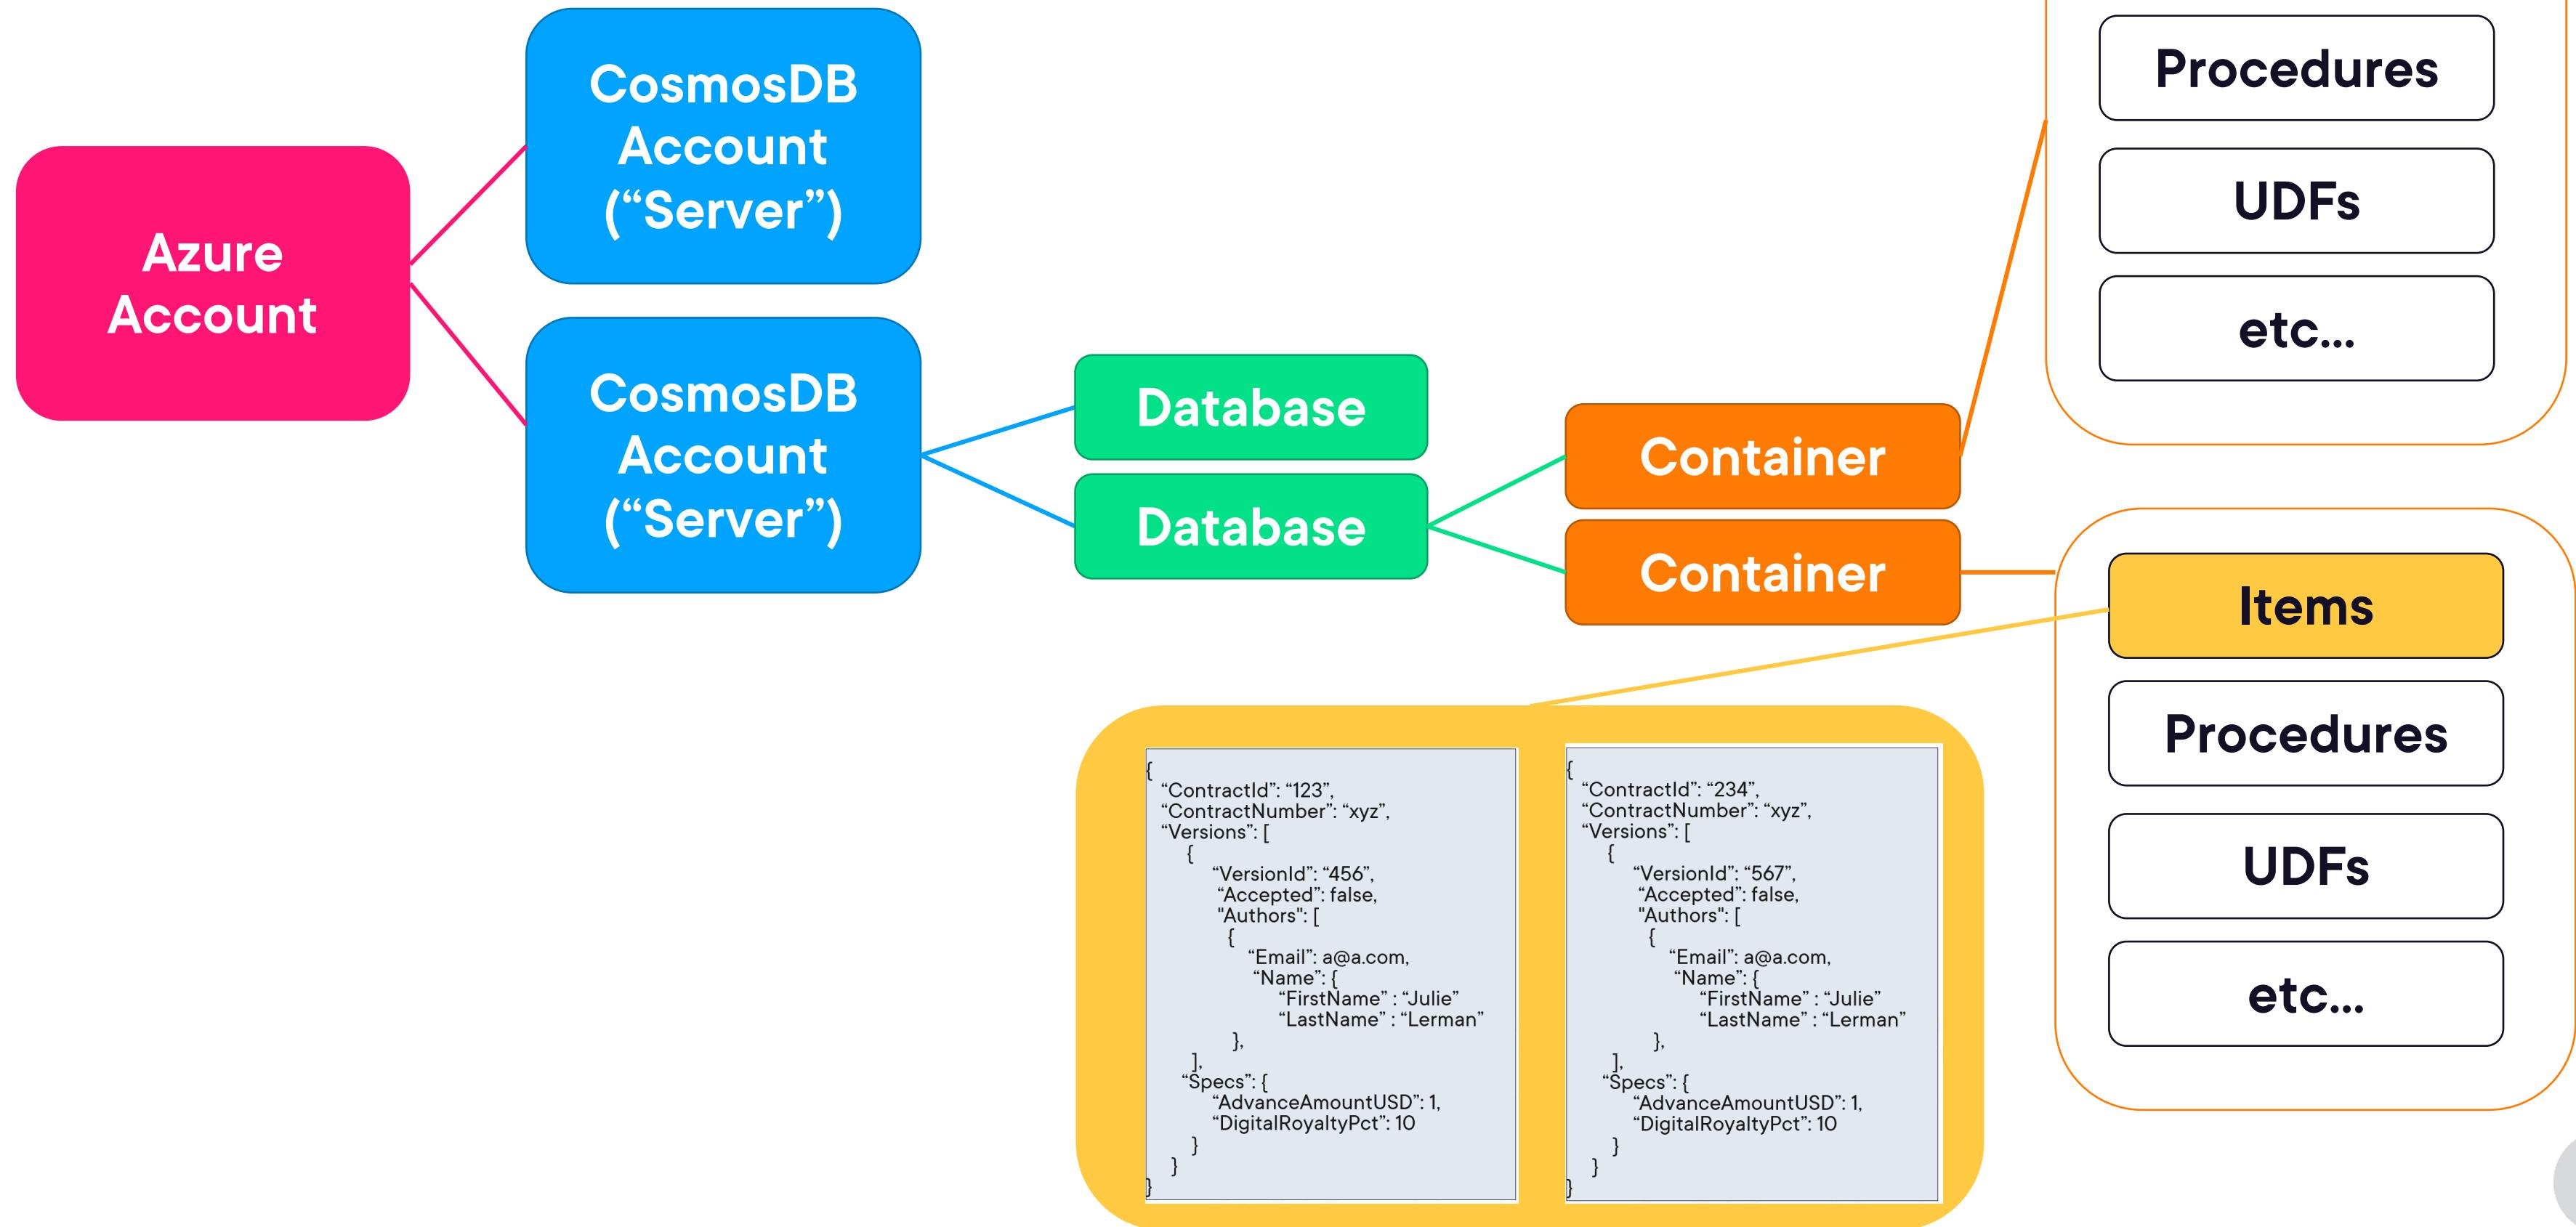
Notes about EF Core 7



Changing the Provider and Removing Incompatible Mappings



CosmosDB Structure





A Global Publishing Powerhouse!

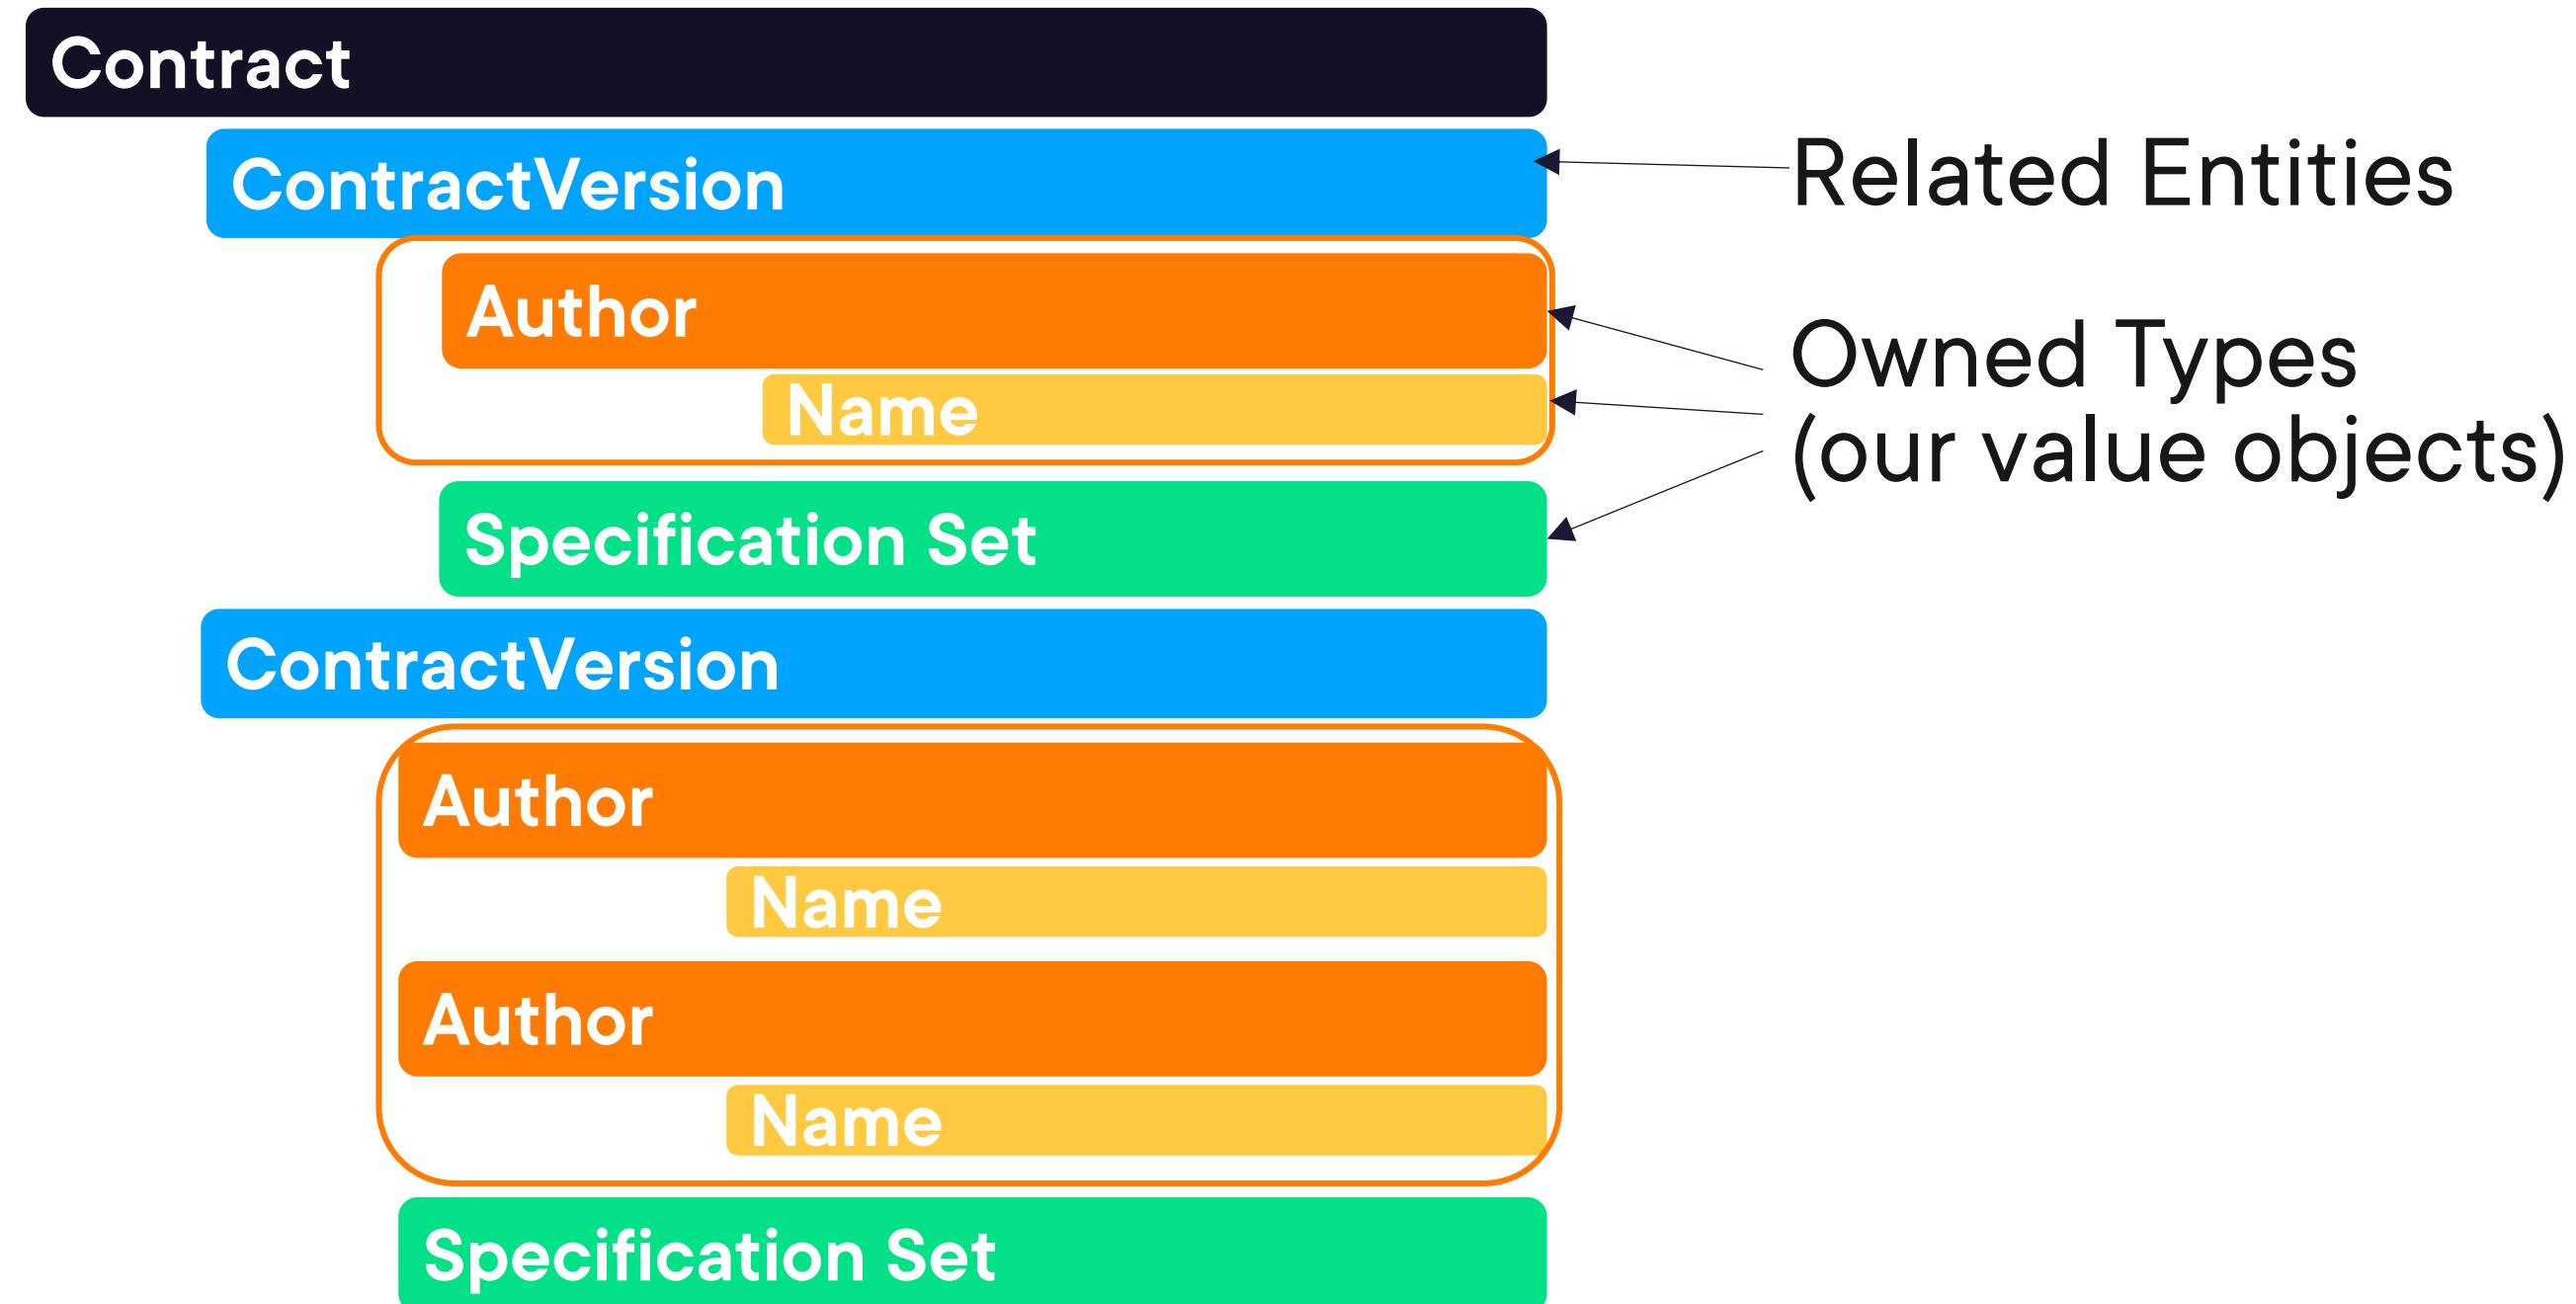
It could become necessary to consider
different mappings to improve throughput and
distribution



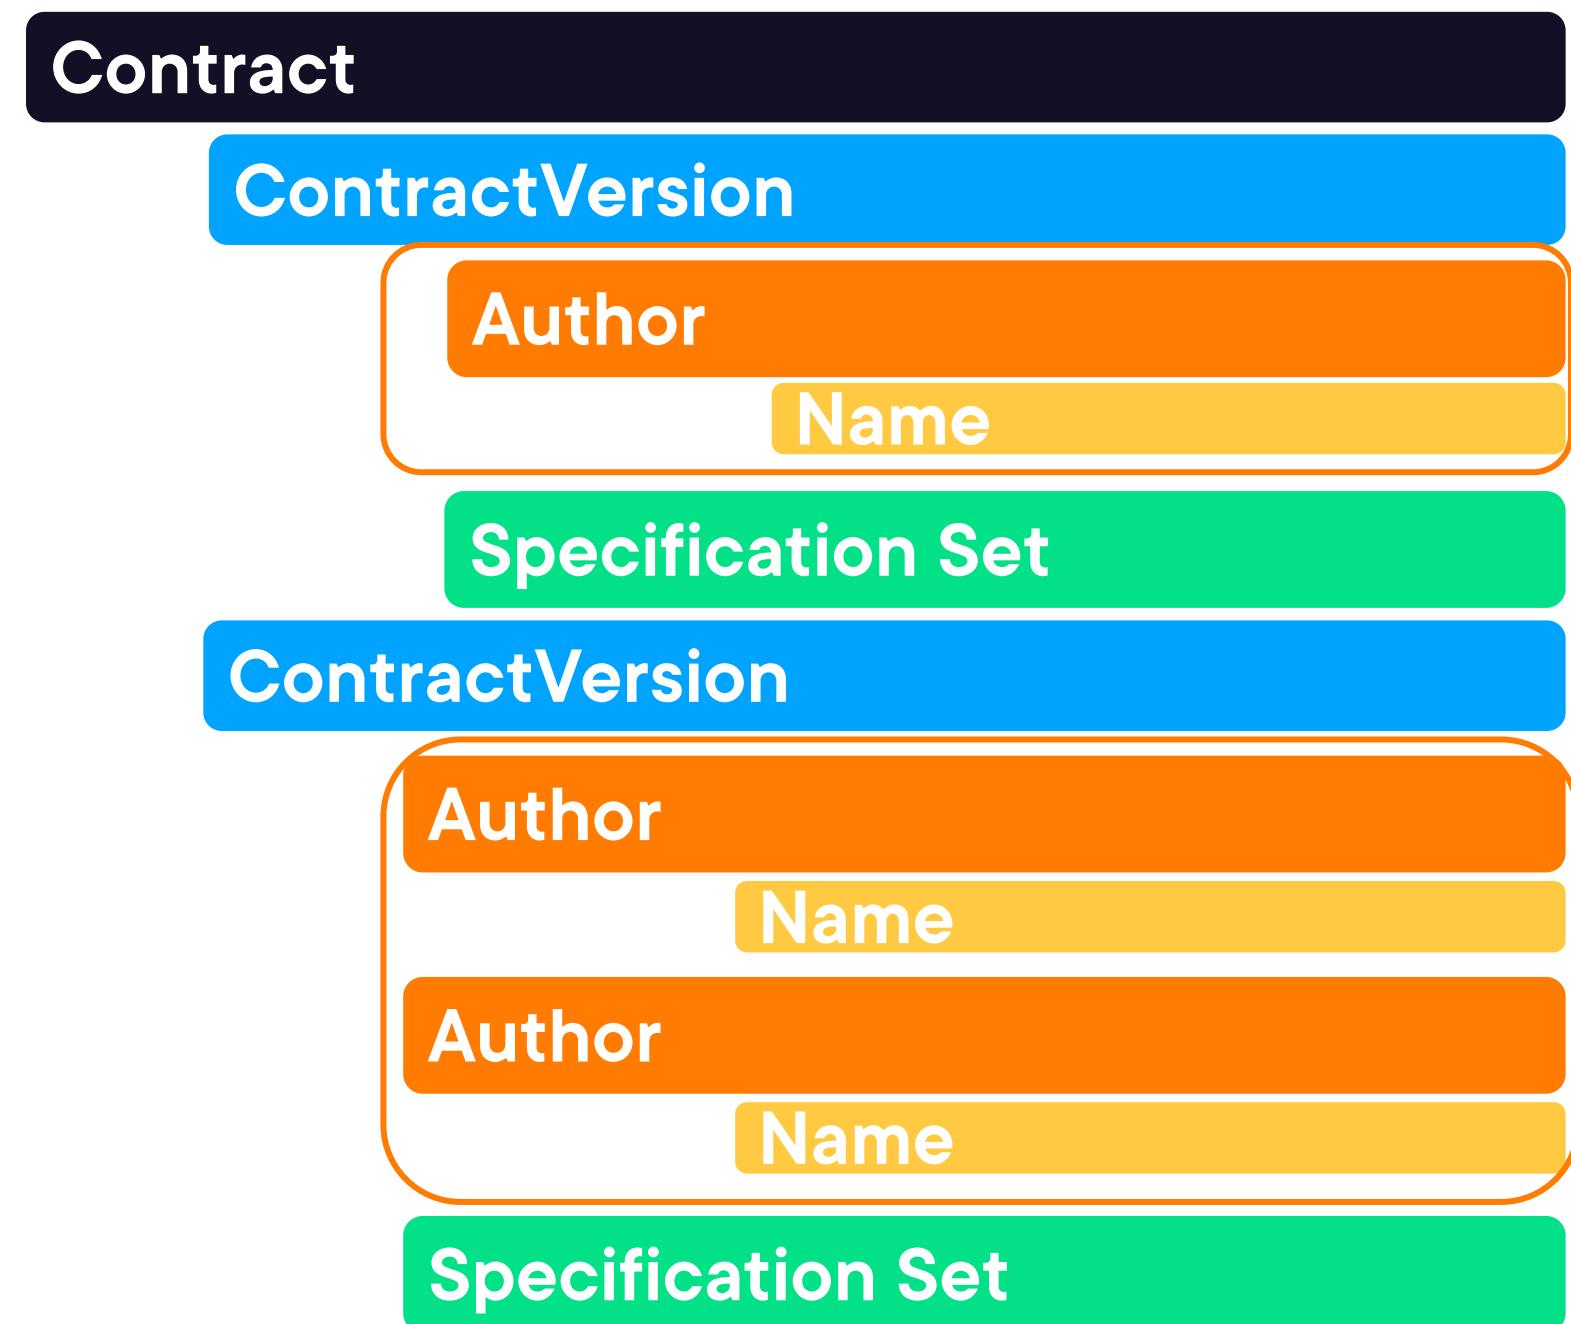
Embedding Value Objects and Related Data in JSON Documents



Default Embedding aka Nesting



Default Embedding aka Nesting



```
{  
  "ContractId": "123",  
  "ContractNumber": "xyz",  
  "Versions": [  
    {  
      "VersionId": "456",  
      "Accepted": false,  
      "Authors": [  
        {  
          "Email": "a@a.com",  
          "Name": {  
            "FirstName": "Julie"  
            "LastName": "Lerman"  
          },  
        ],  
        "Specs": {  
          "AdvanceAmountUSD": 1,  
          "DigitalRoyaltyPct": 10  
        }  
      }  
    }  
  ]  
}
```



Embedded By Convention

Value objects
(EF Core owned types)

- No key property
- Used as a property of another class



Embedded By Convention

**Value objects
(EF Core owned types)**

**Dependent entities
with no DbSet**



Using Integration Tests to Validate the CosmosDB Mapping



Embedded By Convention

Value objects
(EF Core owned types)

Dependent entities
with no DbSet

...or Entity mappings





Fluent API to the Rescue

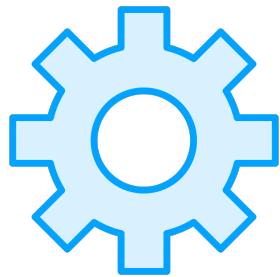
**Apply the needed mappings and still
have the contract versions
embedded**



**There is very helpful info in
EF Core's exception
messages!**



Explicitly Configure the Entity as an Owned Type



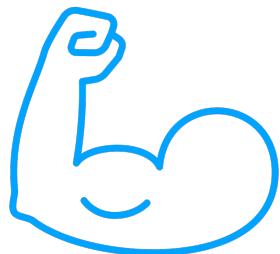
Configure the dependent as an owned entity or owned collection

```
modelBuilder.Entity<Contract>().OwnsMany(c=>c.Versions);
```



Remove (Ignore) the dependent's default mapping, first

```
modelBuilder.Entity<Contract>().Ignore(c => c.Versions);
```



OwnsOne and OwnsMany have 17 overloads!



Single Mapping

```
modelBuilder.Entity<Contract>(c =>
    c.OwnsMany(c => c.Versions,
        owned => owned.Property("_hasRevisedSpecSet")));
```

Multiple Mappings

```
modelBuilder.Entity<Contract>(c =>
    c.OwnsMany(c => c.Versions,
        owned => {
            owned.Property("_hasRevisedSpecSet");
            owned.Property(v => v.Id).ValueGeneratedNever();
        }));
});
```

Configuring Owned Types

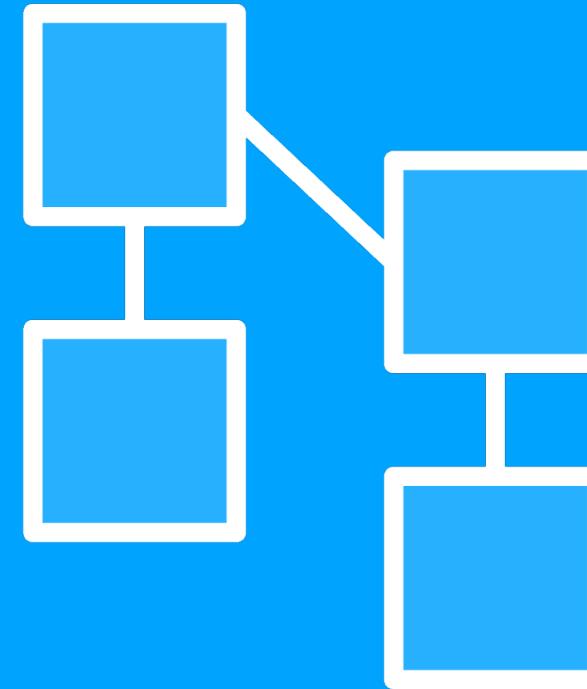
Lambda expression for a single mapping

Lambda method for multiple mappings



Setting Expectations for Querying Related Data





Embedded Objects are Eager Loaded by Default

No need to use `Include` methods



Value Objects and Embedded Dependents

It's standard to retrieve the owned types that are properties of an object

The embedded related objects are treated like owned types and will also come along by default



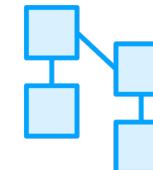
Querying the Root Will Bring Along the Full Aggregate

```
{  
  "ContractId": "123",  
  "ContractNumber": "xyz",  
  "Versions": [  
    {  
      "VersionId": "456",  
      "Accepted": false,  
      "Authors": [  
        {  
          "Email": "a@a.com",  
          "Name": {  
            "FirstName": "Julie"  
            "LastName": "Lerman"  
          },  
        ],  
        "Specs": {  
          "AdvanceAmountUSD": 1,  
          "DigitalRoyaltyPct": 10  
        }  
      }  
    }  
  ]  
}
```

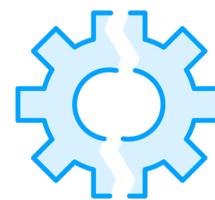
**An aggregate should always
be in a valid state;
therefore, it should be
complete.**



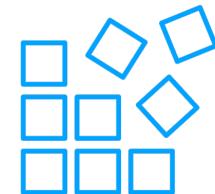
Some of the Current Querying Limitations



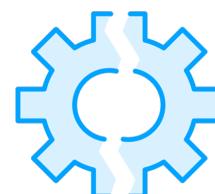
Embedded objects are *always included*, even if you don't want them



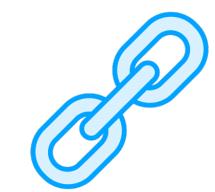
Filtered Include of embedded objects is **not supported**



Projecting the embedded object without the root must be via
NoTracking queries



Filtering or aggregating embedded objects in projections is
not supported



See the Azure Cosmos Task List
(<https://github.com/dotnet/efcore/issues/12086>)



Cosmos Provider Task List

github.com/dotnet/efcore/issues/12086

EF Core Team



Adjusting the Final Tests for Cosmos



EF Core Under the Covers Tricks

Relational Mapping

Treats owned entities as actual entities

Creates shadow properties for fake keys

vs.

Cosmos Mapping

Treats dependent embedded entities like owned entities

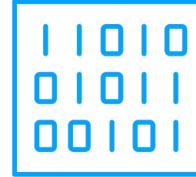




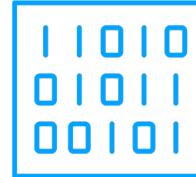
Effortlessly Storing Dictionaries and Lists of Primitives



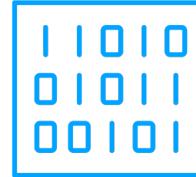
.NET Primitive Types (Collections Map by Convention)



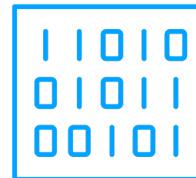
Byte, Sbyte



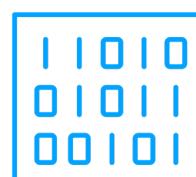
Int16, UInt16, Int32, UInt32, Int64, UInt64, IntPtr, UIntPtr



Char



Double



Single



Cosmos Persists by Convention:

```
public Dictionary<int, string>
    SomePairs{get; set;}
```

```
“SomePairs” : [
    “a” : “the first letter”,
    “z” : “the last letter”
],
```

Dictionaries of <string,string>

```
public List<string> SomeStrings{get; set;}
```

```
“SomeStrings” : [
    “item 1”,
    “item 2”
],
```

Collections of Primitive Types



“We should track social media for potential authors during contracting.”





Azure Cosmos DB

Build or modernize scalable, high-performance apps.

[Try Azure Cosmos DB free](#)[Create a pay-as-you-go account](#)

Fast, distributed NoSQL and relational database at any scale

Develop high-performance applications of any size or scale with a fully managed and serverless distributed database supporting open-source PostgreSQL, MongoDB, and Apache Cassandra. Get [automatic and instant scalability](#), with [SLA-backed](#) single-digit millisecond reads and writes and 99.999 percent availability for NoSQL data. Deploy and scale applications with distributed PostgreSQL using the latest version, tools, and extensions.



Review



Teams have chosen CosmosDB and devs want to benefit from EF Core skills

DDD tactical design fits nicely in JSON

Conventionally embedded objects

Retrieve entire document

Easily support dictionaries & primitive lists

Limitations on querying partial documents

Different mapping needs than RDBMS



Resources from This Module

Cosmos EF provider modifying JSON values with the "id" in identifier

<https://github.com/dotnet/efcore/issues/28600>

Cosmos Provider Task List

<https://github.com/dotnet/efcore/issues/12086>

Accessing JSON via Cosmos Provider

<https://learn.microsoft.com/en-us/ef/core/providers/cosmos/unstructured-data>



Up Next:

Organizing Persistence Logic to Support DDD Design

